

Lab 2: Memory Allocation for XV6 - Alejandro Vargas (avarg116)

Task No.1: Heap Allocator

```
void
fileinit(void)
{
    initlock(&ftable.lock, "ftable");
    //ftable.file= (struct file *)bd_malloc(sizeof(struct file) * NFILE); //change
}

// Allocate a file structure.
struct file*
filealloc(void)
{
    struct file *f;

    acquire(&ftable.lock);
    f = (struct file *)bd_malloc(sizeof(struct file)); //allocate dynamic-sized memory
    release(&ftable.lock);

    if (f) {
        f->ref = 1;
    }

    return f;
}
```

In the initial phase, the main objective was to implement dynamic allocation through the buddy allocator. Improving xv6's memory allocation, the declaration of file[NFILE] had to be removed. I used a pointer to dynamically allocate file structures. The buddy allocator is used to allocate memory for these file structures, allowing dynamic allocation based on available memory. Within struct file* filealloc(void), I used the code inside the function as reference and modified it. Since the code was changed to use dynamic memory allocation for creating a new file structure, the loop was no longer needed.

```
void
fileclose(struct file *f)
{
    acquire(&ftable.lock);
    if (f->ref < 1)
        panic("fileclose");
    if (--f->ref > 0) {
        release(&ftable.lock);
        return;
    }
    bd_free(f); // Deallocate memory using buddy allocator
    release(&ftable.lock);
}
```

Using bd_malloc for allocating memory, it creates a new instance of a struct file dynamically in the heap memory. The conditional statement checks if the memory allocation was completed. ref is set to 1 or null depending on whether a given struct file object is available for use. Then I ran 'alloctest' which passed all tests. It addresses the limitations of the original xv6 design, which had a static allocation of file structures. I also simplified fileclose because ff is not needed but later learned this was not necessary.

Output:

```
init: starting sh
$ allocetest
filetest: start
filetest: OK
$ QEMU: Terminated
[avarg116@sledge cs179f-fall123]$
```

Task No.2: Lazy Page Allocation

In memory management, lazy allocation is a technique that delays the allocation of a resource until processes need it. It's advantageous because it reduces memory resources and lets processes create large sparse data structures. In xv6, the `sbrk()` system call allocates memory and maps it. Implementing lazy allocation for user-space heap memory, we reserve virtual memory and don't allocate immediately. When a page fault happens, it allocates a page and maps it to the faulting address.

```
uint64
sys_sbrk(void)
{
    int n;
    if(argint(0, &n) < 0)
        return -1;
    uint64 oldz = myproc()->sz; //old process size
    myproc()->sz += n; //increase process's size
    if(n < 0) {
        uvmdealloc(myproc()->pagetable, oldz, myproc()->sz); //deallocate memory
    }

    return oldz;
}
```

sysproc.c: in sysproc.c, we were tasked with the removal of page allocation from the `sbrk(n)` system call. I wrote the line `uint64 oldz = myproc()->sz;` which stores the current value of the process's memory size. Similarly, we were instructed to erase the call to `growproc()`.

```
void uvmunmap(pagetable_t pagetable, uint64 va, uint64 size, int do_free)
{
    uint64 a, last;
    pte_t *pte;
    uint64 pa;

    a = PGROUNDDOWN(va);
    last = PGROUNDDOWN(va + size - 1);
    for(;;){
        if((pte = walk(pagetable, a, 0)) == 0) {
            //panic("uvmunmap: walk");
            if(a == last)
                break;
            a += PGSIZE;
            continue;
        }
    }
}
```

vm.c: I made changes to `uvmunmap()` so that it won't panic if all pages aren't mapped. In general, I commented out panics and wrote `if(a == last) break; a += PGSIZE; continue;` to avoid kernel crashes.

Additional changes included handling negative `sbrk()` arguments. This was accomplished through editing the line deallocating memory in the virtual memory system for the current process and return.

```
} else if (r_scause() == 13 || r_scause() == 15) { //check page fault
    uint64 faulting_address = r_stval();
    //kill process
    if(faulting_address > myproc()->sz) {
        myproc()->killed = 1;
    }
    if (faulting_address < myproc()->tf->sp) {
        myproc()->killed = 1;
    }

    uint64 page_boundary = PGROUNDDOWN(faulting_address); //calculate page boundary

    char *mem = kalloc(); //allocate new page
    if (mem == 0) {
        p->killed = 1;
    } else {
        memset(mem, 0, PGSIZE);
        //map new page to page table
        if (mappages(p->pagetable, page_boundary, PGSIZE, (uint64)mem, PTE_W | PTE_X | PTE_R | PTE_U) != 0) {
            kfree(mem);
            p->killed = 1;
        }
    }
}
//edit
```

trap.c: I started by including `else if (r_scause() == 13 || r_scause() == 15)` to check whether a fault is a page fault in `usertrap`. `faulting_address = r_stval();` is extracting faulting virtual addresses using `r_stval()`. Using `kalloc()`, the code proceeds to allocate a new page of physical memory. The instructions suggested stealing code from `uvmalloc()` in `vm.c`. `mappages` deallocates the memory if `mappages` fails. Next, I used `PGROUNDDOWN` to calculate and store the page boundary for a given faulting address.

Usertests: The goal was to make these modifications in the kernel code so that it passes both `lazyttests` and `usertests`. I was not able to complete `usertests` in time, only passing some tests. I got stuck on `pbug` and tried to fix it in `copyout`, `copyin`, and `copyinstr`, but it results in panic: `kerneltrap`.

Testing

```
$ lazytests
lazytests starting
running test lazy alloc
test lazy alloc: OK
running test lazy unmap
test lazy unmap: OK
running test out of memory
test out of memory: OK
ALL TESTS PASSED
$ usertests
usertests starting
test reparent2: OK
test pgbug: scause 0x000000000000000d
sepc=0x00000000000001592 stval=0xeaeb0b5b87f57f5e
panic: kerneltrap
```

```
Task1 grading: allocatest (50/100):
$ make qemu-gdb
OK (5.9s)
Task2 grading: lazytests (25/100):
$ make qemu-gdb
(9.5s)
  lazy: map: OK
  lazy: unmap: OK
Task2 grading: usertests (25/100):
$ make qemu-gdb
Timeout! (150.1s)
  usertests: pgbug: FAIL
    Failed pgbug
  usertests: sbrkbugs: FAIL
    Failed sbrkbugs
  usertests: argptest: FAIL
    Failed argptest
  usertests: sbrkmuch: FAIL
    Failed sbrkmuch
  usertests: sbrkfail: FAIL
    Failed sbrkfail
  usertests: sbrkarg: FAIL
    Failed sbrkarg
  usertests: stacktest: FAIL
    Failed stacktest
  usertests: all tests: FAIL
  ...
    hart 2 starting
    init: starting sh
    $ usertests
    usertests starting
    test reparent2: qemu-system-riscv64: terminating on signal 15 from pid 3705 (make)
MISSING '^ALL TESTS PASSED$'
Score: 75/100
make: *** [grade] Error 1
o [avarg116@sledge cs179f-fall123]$
```